Intro to R and RStudio

SSQL Workshop

Greg Amusu
Swarthmore College

October 23, 2026

Plan for the next two hours

When	What
0:00	**Install together:** R, RStudio, tidyverse, and a Quarto check
0:20	A short tour of RStudio and Quarto, plus good habits
0:40	**Guided walkthrough:** vectors, tibbles, and pipes
1:05	**You work through the handout** while I circulate
1:50	Wrap-up: staying current and where to go next

Install first. Talk second.

We start with installation so anything that goes wrong has time to get fixed.

Part 1

Let’s get set up

1 Get the workshop files

Everything for today is in one place:

github.com/swat-ssql/intro-r-workshop

1. Open the link, click the green **Code** button, then **Download ZIP**
2. **Unzip** it somewhere you'll find it again (not inside Downloads)
3. Keep the page open: the **README** links to the R and RStudio downloads we need next

The two things you're installing

R

The **language and engine**. It does the actual computing.

Free, open source, and maintained by a worldwide community.

RStudio

The **workspace** that makes R easier to use: editor, console, plots, and help in one window.

This is the program you'll open.

Order matters: install R first, then RStudio. Both download links are in the repo's README.

2 Install R

Mac | Windows

1. Go to cran.r-project.org and click **Download R for macOS**
2. Download the installer (`.pkg`) that matches your chip: **Apple silicon** (M1, M2, M3, ...) or **Intel**
   Not sure? Apple menu → About This Mac.
3. Open the `.pkg` and click through the installer
4. Open **Terminal** (⌘ + Space, type “terminal”) and run:

```
xcode-select --install
```

Say yes to installing the command line tools. Already have them? You’re set.

2 Install R

Mac | Windows

1. Go to cran.r-project.org and click **Download R for Windows**
2. Click **base**, then download the installer and run it (the defaults are fine)
3. *Optional:* install **Rtools** from the Rtools page, choosing the version that matches your R. You only need it if a package install asks for it.

3 Install RStudio

1. Go to posit.co/download/rstudio-desktop
2. Download the installer and run it (Mac: drag RStudio into Applications)
3. Double-click `intro-r-workshop.Rproj` to open the project
4. Type this into the **Console** (bottom left) and press Enter:

```
R.version.string
#> [1] "R version 4.6.x (...)"
```

An R version means RStudio and R are talking. The Files pane should list `handout.qmd`.

4 Install the tidyverse

```
install.packages("tidyverse", repos = "https://cloud.r-project.org")
```

- Takes a few minutes. Lots of scrolling text is normal.
- If it asks *“Do you want to install from sources...?”*, answer **No**.
- You install a package **once**. You load it with `library()` **every session**:

```
library(tidyverse)
```

You should see **“Attaching core tidyverse packages”** and a list of packages.

5 Check that Quarto works

You'll need **Quarto** to turn a `.qmd` file into a finished document. **RStudio already includes it**, so there's nothing extra to download.

1. **File → New File → Quarto Document...**, then click **Create** (leave HTML selected). If it offers to install packages, say **yes**
2. Save the file, then click the blue **Render** arrow. An HTML preview should open
3. Now try a PDF: change `format: html` to `format: typst` in the header and **Render** again. Typst comes with Quarto, so nothing else to install

> ⚠ **Render says Quarto isn't found, or the PDF fails?**
>
> Install Quarto from quarto.org/docs/download. HTML always works, and that's fine for today.

Checkpoint

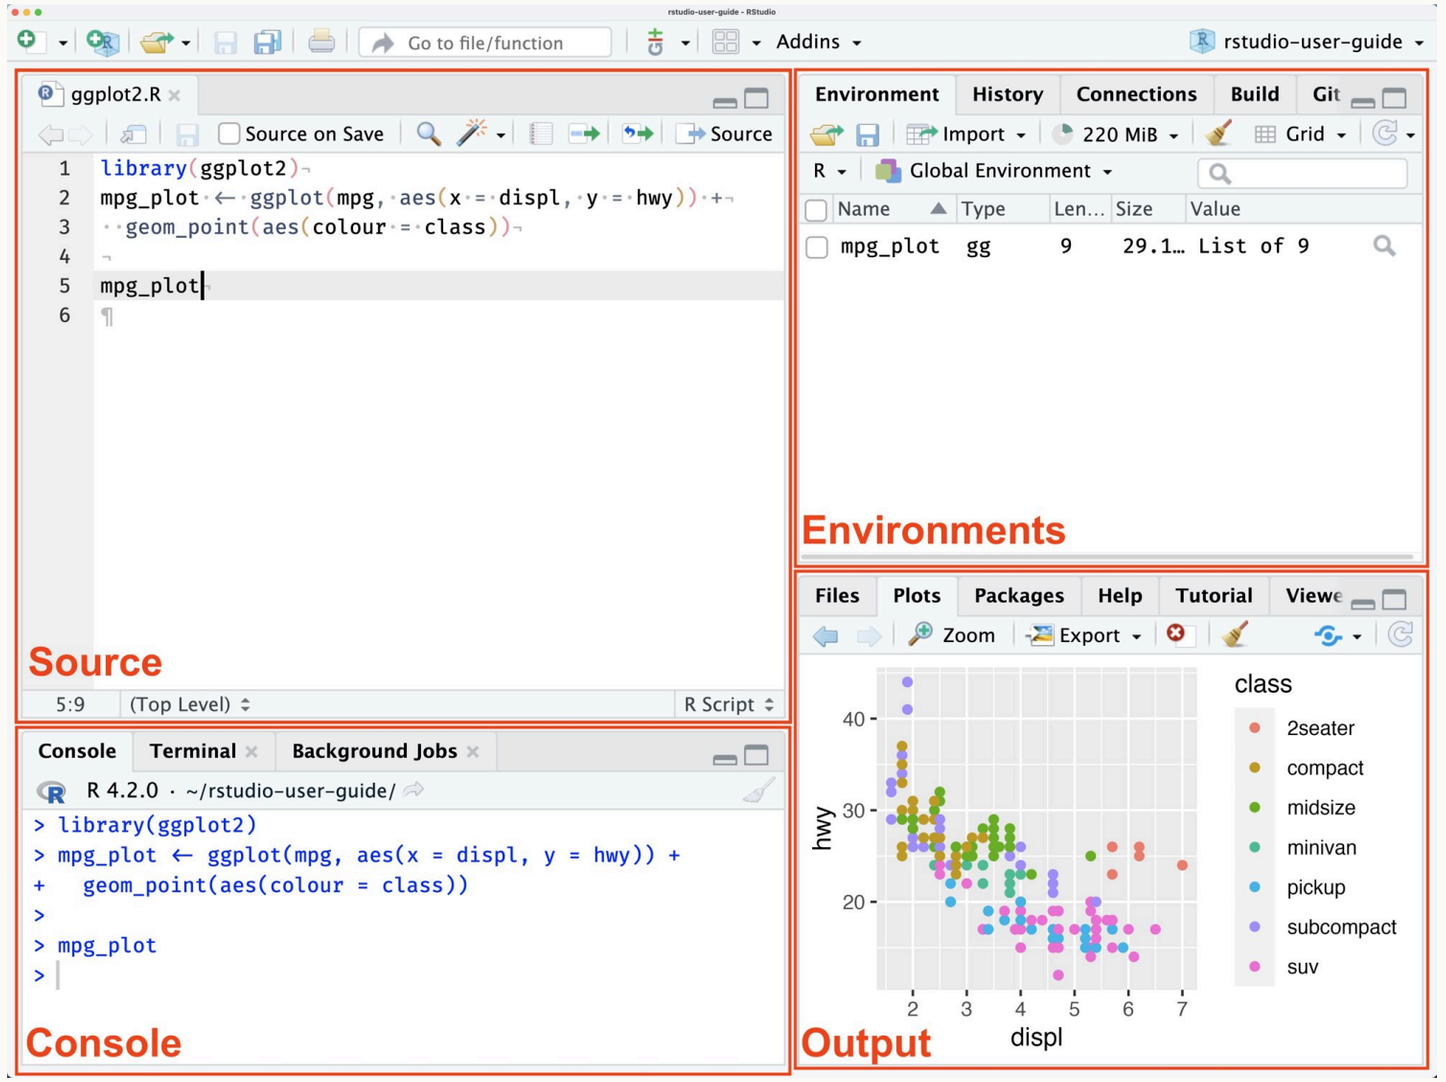
Hand up if any of these aren't green:

- [ ] The workshop project is open (the Files pane shows `handout.qmd`)
- [ ] `R.version.string` prints a version
- [ ] `library(tidyverse)` runs without an error
- [ ] A test Quarto document renders to HTML and to PDF (`format: typst`)

Part 2

A quick tour

A tour of RStudio



Where does my code live?

Console

Try things out. Quick answers.

Not saved.

Script (`.R`)

A saved list of commands.

Run a line with **⌘/Ctrl + Enter**.

Quarto (`.qmd`)

Code + writing + results in one document.

Press **Render** to build it.

Rule of thumb: if you want to keep it, it goes in a file.

Can I only use R in a Quarto file in RStudio?

No. R is the same everywhere. Only the place you type changes.

Where	Good for
R console: the Console pane, or type `R` in a terminal	quick experiments
Script (`.R`)	just code; run a line at a time or the whole file
Quarto (`.qmd`)	code + writing + results in one document
Other editors: Positron, VS Code	same R in a different workspace

```
Rscript handout_code.R     # run a whole script from a terminal
```

See the difference

`handout_code.R` in the repo holds the same code as `handout.qmd`, as a plain script. Open them side by side.

Quarto: code, writing, and results together

````
---
title: "My analysis"
format: html
---

Some **text** written in Markdown.

```{r}
mean(c(2, 4, 6))
```
````

- The top block (**YAML**) sets the title and output format: PDF, HTML, or Word
- Text goes in normal paragraphs. Code goes in **chunks**
- Press **Render** and Quarto runs the code and builds the document
- `handout.qmd` is a Quarto file, and it renders to PDF with Typst

Projects and the working directory

`getwd()` tells you which folder R is looking in right now.

- Open the workshop as a **Project**: double-click `intro-r-workshop.Rproj` (or File → Open Project)
- Now file paths are relative to the project folder:

```
read_csv("data/subset_saf.csv")
```

- Avoid `setwd("/Users/you/Desktop/...")`. It breaks on every other computer.

Objects, functions, packages

```
x <- c(2, 4, 6)         # <- stores a value in an object called x
mean(x)                 # functions look like name(inputs)
#> [1] 4
library(tidyverse)      # packages add more functions
```

- **Assign** with `<-` (shortcut: **Option/Alt + -**)
- **Call** a function with `name(inputs)`
- **Install** a package once, **load** it every session

Part 3

Good habits from day one

Stay organized

1. **One project folder** per course or analysis, with data, code, and output kept in clear places
2. Work in a **script or Quarto file**, not just the console
3. **Raw data stays raw.** Never edit the original file. Do all cleaning in code so you (or anyone else) can rerun it
4. Use clear names: `snake_case`, and `turnout` rather than `df2`. Avoid spaces in file names
5. Comment the **why**, not just the what

Know your data before you analyze it

What is one row?

A county-year? A police stop? A person? This is your **unit of observation**.

Read the codebook

What does each variable mean, in what units, and how are **missing values** coded (`NA`, `-99`, `999`)?

Look before you leap

`glimpse()`, `summary()`, `count()`. Do the ranges make sense? An age of 0 or 999 should stop you.

Remember the people

Behind each row of social data is a person. Handle it with care.

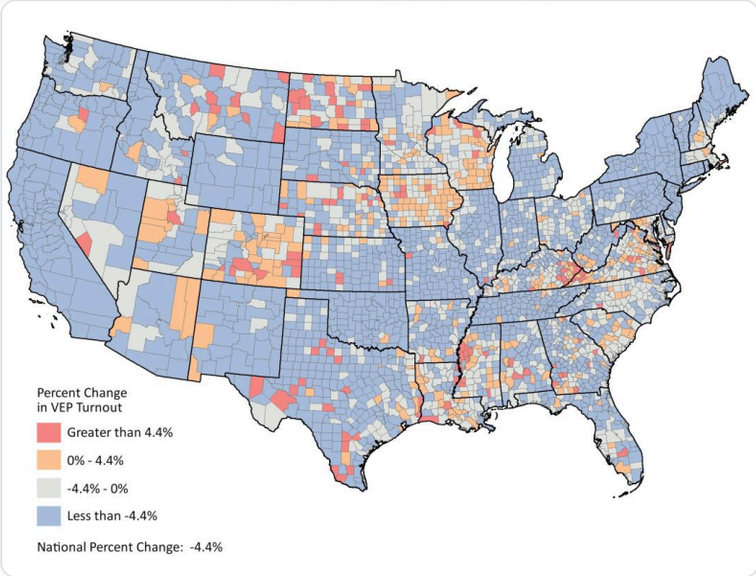
Make it reproducible

- **Restart R and run everything top to bottom** (Session → Restart R, then Render) before you submit. If it works from scratch, it works.
- **Don't save your workspace.** Tools → Global Options → General: set *Save workspace to .RData on exit* to **Never** and untick *Restore .RData into workspace at startup*
- **Know your versions:** `sessionInfo()` lists R and package versions
- Later: version control with **git and GitHub** (this workshop lives in a repo)

Part 4

Guided walkthrough: from vectors to tibbles

The data you'll meet today



Voter turnout, 2008 vs. 2012. Vectors and tibbles.



Disaster aid and the incumbent vote. Healy and Malhotra (2009).



NYC stop-and-frisk. Your questions at the end.

Follow along in the handout

1. Open `handout.qmd` (in the project you opened earlier) and put your name in the header
2. Work **top to bottom**: run each chunk with the green ▶ (or **⌘/Ctrl + Shift + Enter**) and read the output
3. **Run the code as I show it.** The vectors-to-pipes code is already in the handout
4. Some chunks are blank on purpose. `handout_guide.pdf` shows the finished versions
5. Do the **Workshop Questions** on your own at the end
6. **Render** to PDF when you're done

Working with vectors

We'll take the turnout figures from Trende's chart in the handout and **combine** them into vectors, one per column.

```
actual_2008    <- c(97677100, 16564353, 11174365, 2629262, 3418041)
projected_2012 <- c(99139605, 17455952, 13134504, 3300869, 3910201)
actual_2012    <- c(93033145, 16797651, 12921270, 3230318, 2584254)
```

Math works **element by element**:

```
difference <- actual_2012 - projected_2012
difference
#> [1] -6106460  -658301  -213234   -70551 -1325947
```

Picking out elements: `[ ]`

```
actual_2008[2]
#> [1] 16564353

actual_2012[1] - projected_2012[1]
#> [1] -6106460
```

Concept check: what did each line do?

The first picks the **2nd** element. The second subtracts the **1st** projected value from the **1st** actual value.

Common vector functions

```
length(actual_2008)   # 5
mean(actual_2008)     # about 26.3 million
median(actual_2008)   # 11174365
range(actual_2008)    # 2629262 97677100
sum(actual_2008)      # 131463121
```

The handout adds `min()`, `max()`, `prod()`, `log()`, and `unique()`. Try them on your own.

Naming vectors

```
names(actual_2008)    # no names yet
#> NULL

names(actual_2008) <- c("White", "Black", "Hispanic", "Asian", "Other")
actual_2008
#>    White    Black Hispanic    Asian    Other
#> 97677100 16564353 11174365  2629262  3418041
```

```
actual_2008["White"]
#>    White
#> 97677100
```

From vectors to tibbles

A dataset is just **named vectors side by side**.

- Vectors are single columns of numbers or labels
- Real data are **tables**: many columns, each a variable
- **Tibbles** (from the tidyverse) are tables that print more clearly and play nicely with tidyverse tools
- Let’s combine our vectors into one

Building a tibble

```
turnout <- tibble(
  group = c("White", "Black", "Hispanic", "Asian", "Other"),
  actual_2008 = actual_2008,
  projected_2012 = projected_2012,
  actual_2012 = actual_2012
)
turnout
```

```
#> # A tibble: 5 × 4
#>   group    actual_2008 projected_2012 actual_2012
#>   <chr>          <dbl>          <dbl>       <dbl>
#> 1 White       97677100       99139605    93033145
#> 2 Black       16564353       17455952    16797651
#> 3 Hispanic    11174365       13134504    12921270
#> 4 Asian        2629262        3300869     3230318
#> 5 Other        3418041        3910201     2584254
```

Try breaking it

What happens if you delete `, "Other"` so that

```
group = c("White", "Black", "Hispanic", "Asian")
```

R refuses, and tells you why:

```
#> Error in `tibble()`:
#> ! Tibble columns must have compatible sizes.
#> • Size 4: Existing data.
#> • Size 5: Column `actual_2008`.
```

Columns in a table must all be the **same length**. Reading error messages is a skill.

Pipes: `%>%` and `|>`

- A pipe passes the result of one step into the next, so code reads **left to right**
- `%>%` (tidyverse) and `|>` (base R) do the same job here
- You can **chain** steps instead of nesting them

```
mean(log(x))                  # nested: read from the inside out

x %>% log() %>% mean()        # piped: read left to right
x |> log() |> mean()          # same thing, base R pipe
```

Pulling out and adding columns

`$` pulls out one column **by name**: `object$column`

```
turnout$actual_2012 - turnout$projected_2012
#> [1] -6106460  -658301  -213234   -70551 -1325947
```

`mutate()` **adds** a column to the tibble:

```
turnout <- turnout %>%
  mutate(diff_2012 = actual_2012 - projected_2012)
```

Just want the vector back? Add `%>% pull(diff_2012)`. The handout shows how.

Adding a column: the result

```
turnout
#> # A tibble: 5 × 5
#>   group    actual_2008 projected_2012 actual_2012 diff_2012
#>   <chr>          <dbl>          <dbl>       <dbl>     <dbl>
#> 1 White       97677100       99139605    93033145  -6106460
#> 2 Black       16564353       17455952    16797651   -658301
#> 3 Hispanic    11174365       13134504    12921270   -213234
#> 4 Asian        2629262        3300869     3230318    -70551
#> 5 Other        3418041        3910201     2584254   -1325947
```

The new `diff_2012` column sits at the end.

Which group had the biggest drop?

`arrange()` sorts the rows. `slice(1)` keeps the first one.

```
turnout %>% arrange(diff_2012) %>% slice(1)
```

```
#> # A tibble: 1 × 5
#>   group actual_2008 projected_2012 actual_2012 diff_2012
#>   <chr>       <dbl>          <dbl>       <dbl>     <dbl>
#> 1 White    97677100       99139605    93033145  -6106460
```

Does the answer surprise you? Why or why not? What other new columns would help here?

Part 5

Your turn

Getting unstuck

1. **Built-in help:** `?mean` or `??regression`. The examples at the bottom are the best part.
2. **Read the error message.** Common causes: a typo, a missing `)` or quote, forgetting `library()`, a wrong file path.
3. **Search the exact message** along with “R” or “tidyverse”.
4. **Good references:** R for Data Science and tidyverse.org
5. **Ask.** Me, a neighbor, or your instructor.

> **AI chat tools**
>
> They can explain error messages, but always run what they give you and check it does what you think. Follow your course’s policy.

Your turn!

Open `handout.qmd` and begin.

When you're stuck:

- Check the **Console** for an error message and read it
- Check `getwd()`: is R in the right folder?
- Ask a neighbor, then ask me

Workshop files: github.com/swat-ssql/intro-r-workshop

Part 6

Wrap-up

Keeping R, RStudio, and Quarto up to date

	How often	How to update
R	A new version each spring (4.6.0 came out April 2026), plus small fixes	Download and install the new version from CRAN, then **reinstall your packages**
RStudio	Roughly monthly. Numbered by year.month (e.g., 2026.10)	Help → Check for Updates
Quarto	Comes **bundled with RStudio**, so it updates when RStudio does	Update RStudio, or download from quarto.org
Packages	Whenever their authors update them	Tools → Check for Package Updates

Don't update in the middle of an assignment.

Update at the start of a term or a break. If everything works, you can leave it alone.

RStudio isn’t the only option: meet Positron

- A newer editor from **Posit**, the makers of RStudio, built on the same open-source base as **VS Code**
- Works with **R and Python** side by side
- Same R, same packages, same Quarto, so what you learn today carries over
- Posit says **RStudio is not going away** and will keep being maintained
- For this course we’ll stick with **RStudio**, so screenshots and instructions match

Curious? positron.posit.co

Where to go next

- Keep your project folder. You'll reuse these skills all semester
- Read: R for Data Science (2e), the first few chapters
- Try: **Positron**, if you're curious about Python or VS Code
- Keep asking questions: nobody learns R without getting stuck first

Keep learning with the SSQL

Upcoming R workshops. Register on the SSQL website.

Date	Workshop
Oct. 28	U.S. Midterm Election Modeling
Nov. 5, 6	Exploring Data in R
Nov. 9, 13	Data Visualization in R

Need help with a course or project?

Schedule an **office hours appointment** at the SSQL for course or project support.

Thank you!